

Introduction to Akka Streams

Reactive Stream Processing with Back Pressure



Outline

- Reactive Streams
 - Motivation
 - Concepts
 - Back-Pressure
 - Use Cases
- Akka Streams
 - Materializing and running
 - Stream completion
 - Example
 - Sources
 - Sinks
 - Flows
 - Combinators
 - File I/O
 - Akka Http
 - Server
 - Client
 - GraphDSL

Reactive Streams

Motivation

Goal:

Processing of data streams ...

Challenges:

- ... of unknown length?
- ... with memory constraints?
- ... with flow capacity constraints?
- ... non-blocking / asynchronously?
- “Big Data”?

Motivation (2)

Solution:

- Standard API for stream processing
- Flow control through “Back-Pressure”
- Implementation with Akka actors
- Functional composition of processing stages
- Bounded buffers of processing elements
- Reusability of defined flows

Concepts

Stream:

Active process that involves moving and transforming data from “upstream” to “downstream” (= materialized graph)

Graph:

Connected pathway in directed graph, defining the topology/blueprint of a stream

Graph Stage:

Building block of graph,

Concepts (2)

Source:

Graph element with 1 output, emitting elements into the stream

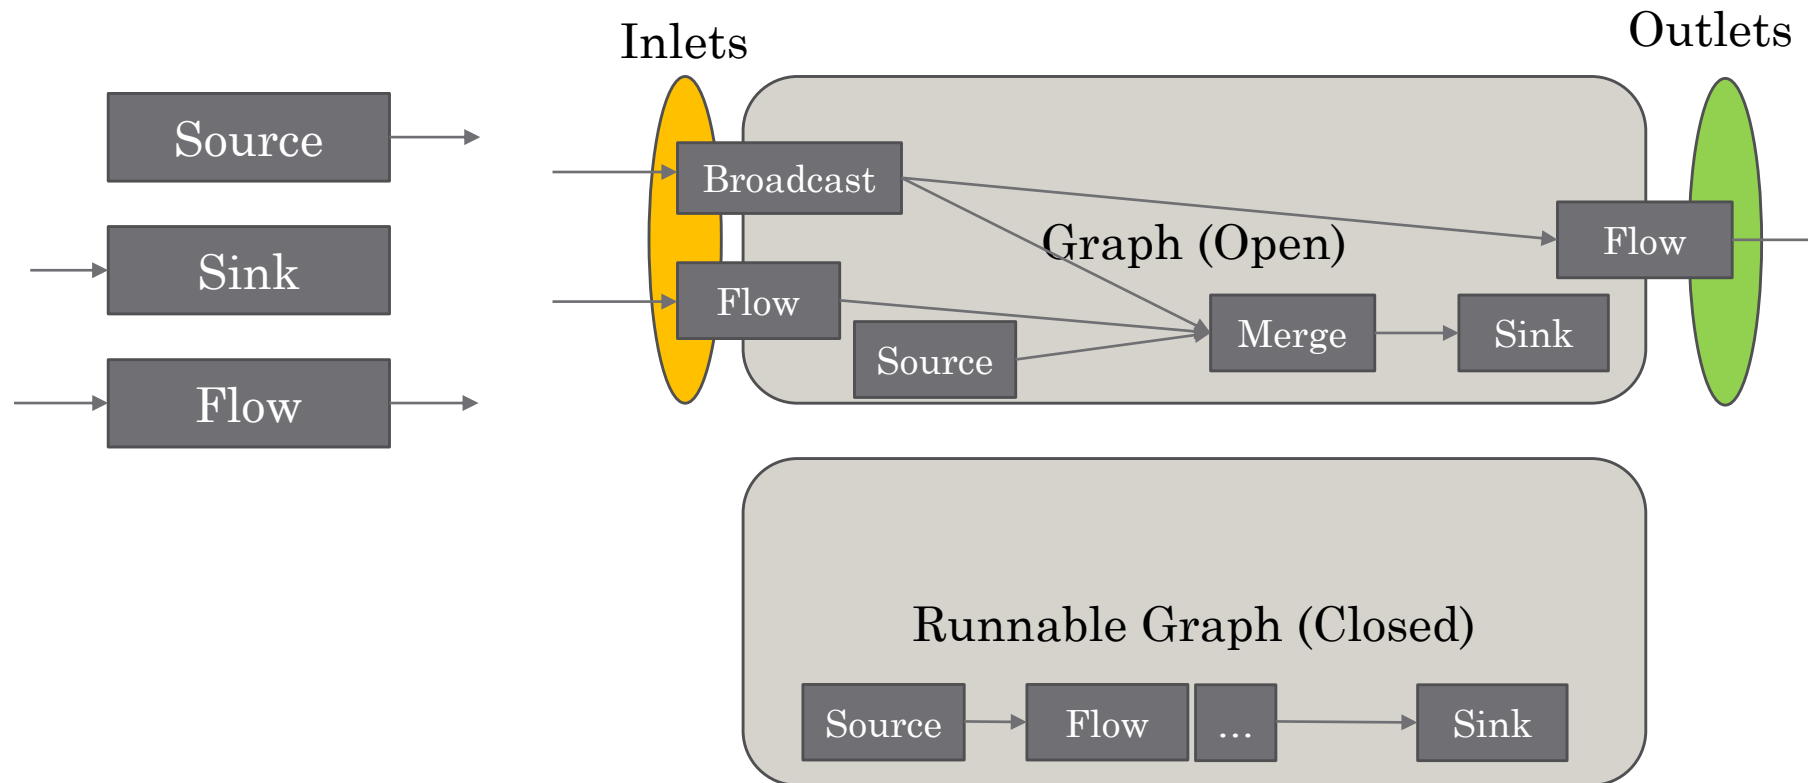
Sink:

Graph element with 1 input, consuming elements from the stream

Flow:

Graph element with 1 input and 1 output

Concepts (3)



Back-Pressure

Problem:

Fast Publisher, Slow Subscriber:

Publisher emits more elements into a stream than the Subscriber can handle

→ Buffer overflows, elements lost (bounded)

→ Out of Memory Errors (unbounded)

Solution:

Subscriber signals the Publisher the rate of elements it can consume

Use Cases

- I/O Streaming
 - Reading/writing large files
 - Processing/transforming data from files while reading
 - Streaming data from/to sockets
- Handling of Requests (as Flow)
 - HTTP requests (Akka Http)
 - WebSockets
- Complex stream processing by composition of processing stages
 - Custom Graphs/GraphStages
 - “GraphDSL”

Akka Streams

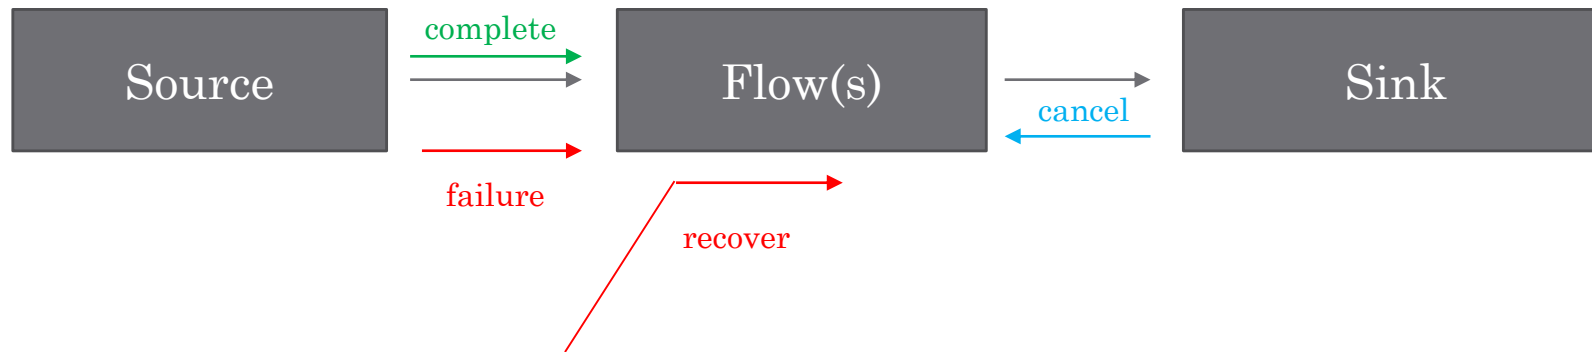


Materializing and running

- Streams need to be materialized from graphs (blueprint → stream)
- Akka Actor System, Actor Materializer required
- Stream processing happens in dedicated threadpool
- Each graph stage gets materialized to runtime value, to access the status **while the stream is running**
 - Usually `Future[_]` when graph stage contains
 - Value “NotUsed” for graph stages where the materialized is irrelevant (e.g. `map()`)
 - Combining
 - With `map()`, `filter()`, `via()` ... the materialized value of the source is kept by default
 - With `runWith()` the materialized value of the sink is kept by default

Graph completion

- A stream finishes:
 - when the source signals no more elements to downstream (“complete”)
 - when the sink requests no more elements from upstream (“cancel”)
 - when the stream is failed (e.g. Exception in flow) and no recovery is done



Example

```
object SimpleExample extends App {  
  
  implicit val system: ActorSystem = ActorSystem("example")  
  implicit val materializer: Materializer = ActorMaterializer()  
  
  val source = Source(List("apple", "pear", "orange"))  
  
  val sink = Sink.ignore  
  
  // Materialize and run stream  
  val streamResult = source.runWith(sink)  
  
  // Handle future  
  Await.ready(streamResult, atMost = 10 seconds)  
  
}
```

Example: Sources

```
class Sources {  
  
    // Empty  
    val sourceEmpty = Source.empty[String]  
  
    // From Iterable  
    val sourceIterable = Source(1 to 1000)  
    val sourceIterable2 = Source(List(1, 2, 3, 99))  
    val sourceIterableInfinite = Source(Stream.from(1))  
  
    // From File  
    val sourceFile = FileIO.fromPath(Paths.get("/my/file"))  
  
    // Repeat same value  
    val sourceRepeat = Source.repeat("Hello World")  
  
    // From Future  
    val promise = Promise[Int]()  
    val sourceFuture = Source.fromFuture(promise.future)  
    promise.success(1337)  
  
}
```

Example: Sinks

```
class Sinks {  
  
  // Ignore values (drain stream)  
  val sinkIgnore = Sink.ignore  
  
  // head, headOption, last, lastOption  
  val sinkHead = Sink.head[String]  
  val sinkHeadOption = Sink.headOption[Int]  
  val sinkLast = Sink.last[String]  
  val sinkLastOption = Sink.lastOption[String]  
  
  // fold (e.g. count chars)  
  val sinkFoldCount = Sink.fold[Int, String](zero = 0)((numChars, nextString) => numChars + nextString.length)  
  val sinkFoldConcat = Sink.fold[ByteString, ByteString](ByteString.empty)(_ ++ _)  
  
  // collect all values  
  val sinkSeq = Sink.seq[String]  
}
```


Example: Flows

```
class Flows {  
  
  Source(1 to 100000)  
    .filter(_ % 2 == 0 /* Even numbers only */)   
    .drop(15 /* drop first 15 elements */)   
    .map(nr => ByteString(s"$nr Line $nr\n")) /* Convert each to ByteString as "[33] Line 33\n" */)   
    .to(FileIO.toPath(Paths.get("out.txt"))) /* Write to file */   
    .run()  
  
  val fruit = Source(List("apple", "banana", "orange"))  
  val prices = Source(List("1,33", "2,44", "0,99")).map(_ + "€")  
  
  fruit  
    .zipWith(prices) ("The " + _ + " costs " + _)   
    .runForeach(println)  
  // The apple costs 1,33€  
  // The banana costs 2,44€  
  // The orange costs 0,99€  
  
}
```

Combinators

Merge:

Combine multiple sources to a single, combined source (Fan-in)

Broadcast:

Emit to multiple sinks (Fan-out)

Zip:

Combine each 1st, 2nd, ... elements from multiple streams

Concat:

Pull from 1st stream, and then from 2nd stream after the first finished

FileIO

```
class FileIOExample {  
  
    // File source: Read data (ByteString) from file  
    val fileSource = FileIO.fromPath(Paths.get("in.dat"), chunkSize = 8192)  
  
    // File sink: Store data (ByteString) into file  
    val fileSink = FileIO.toPath(Paths.get("out.dat"))  
  
    // Stream from one file to another  
    fileSource.runWith(fileSink)  
  
}
```

Akka Http: Server

```
val route: Route = {
  path("url") {
    get {
      complete {
        val source = FileIO.fromPath(Paths.get("myfile.txt"))
        // Stream download from file or other source
        HttpEntity(ContentTypes.`application/octet-stream`, source)
      }
    } ~
    post {
      extractRequestEntity { entity =>
        val sink = FileIO.toPath(Paths.get("myfile.txt"))
        // Stream upload into file or other sink
        onSuccess(entity.dataBytes.runWith(sink)) { streamResult =>
          complete {
            if (streamResult.wasSuccessful) {
              "upload complete"
            } else {
              StatusCodes.InternalServerError -> s"error storing file: ${streamResult.getError.getMessage}"
            }
          }
        }
      }
    }
  }
}
```

Akka Http: Client

```
val httpFlow: Flow[HttpRequest, HttpResponse, Any] = Http().outgoingConnectionHttps("illucit.com", port = 443)

// Single Request

val request = Get("https://www.illucit.com/impressum/")

val singleResponse: Future[String] =
  Source.single(request)
    .map(_ => addHeader(Cookie("qtrans_front_language" -> "de")))
    .via(httpFlow)
    .mapAsync(1) (Unmarshal(_).to[String])
    .runWith(Sink.head)

// Flow

val responseFlow: Flow[Uri, String, NotUsed] = Flow[Uri]
  .map(uri => Get(uri))
  .map(_ => addHeader(Cookie("qtrans_front_language" -> "de")))
  .via(httpFlow)
  .mapAsync(1) (Unmarshal(_).to[String])

val responses: Future[Seq[String]] =
  Source(List(Uri("https://www.illucit.com/"), Uri("https://www.illucit.com/robots.txt"), Uri("https://www.illucit.com/sitemap_index.xml")))
    .via(responseFlow)
    .runWith(Sink.seq)
```

GraphDSL

- Domain specific language to defined graphs from basic building blocks
- Inner graph stages define behavior of graph
- Non-connected graph stages define the “shape of the graph”
 - Closed Shape → Graph has no unconnected ports → RunnableGraph
 - FlowShape: One Input and One Output can be connected
 - SourceShape: One Output
 - SinkShape: One Input

GraphDSL: Example

```
RunnableGraph.fromGraph(GraphDSL.create() {  
  implicit builder: GraphDSL.Builder[NotUsed] =>  
    import GraphDSL.Implicits._  
  
    val in = Source(1 to 10)  
    val out = Sink.ignore  
  
    val bcast = builder.add(Broadcast[Int](2))  
    val merge = builder.add(Merge[Int](2))  
  
    val f1, f2, f3, f4 = Flow[Int].map(_ + 10)  
  
    in ~> f1 ~> bcast ~> f2 ~> merge ~> f3 ~> out  
    bcast ~> f4 ~> merge  
    ClosedShape  
})
```

GraphDSL: Example (2)

```
Flow.fromGraph(GraphDSL.create() {  
  implicit builder: GraphDSL.Builder[NotUsed] =>  
    import GraphDSL.Implicits._  
  
  val in = builder.add(Flow[Int])  
  val out = builder.add(Flow[String])  
  
  val bcast = builder.add(Broadcast[Int](2))  
  val merge = builder.add(Merge[Int](2))  
  
  val f1, f2, f3, f4 = Flow[Int].map(_ + 10)  
  
  val toStr = Flow[Int].map(_.toString)  
  
  in ~> f1 ~> bcast ~> f2 ~> merge ~> f3 ~> toStr ~> out  
  bcast ~> f4 ~> merge  
  FlowShape(in.in, out.out)  
})
```




Thank you for your attention !

github.com/metaxmx

simon@illucit.com

www.illucit.com

